

# Sequential data

- ① Time series
  - ▶ Financial data analysis: stock market, commodities, Forex
  - ▶ Healthcare: pulse rate, sugar level (from medical equipment and wearables)
- ② Text and speech: speech understanding, text generation
- ③ Spatiotemporal data
  - ▶ Self-driving and object tracking
  - ▶ Plate tectonic activity
- ④ Physics: jet identification
- ⑤ etc.

# Sequence modelling I

## Sequence classification

- ①  $\mathbf{x} = x_1, x_2, \dots, x_n, x_i \in V$  - objects
- ②  $y \in \{1, \dots, L\}$  - labels
- ③  $\{(\mathbf{x}^{(1)}, y_1), (\mathbf{x}^{(2)}, y_2), \dots, (\mathbf{x}^{(m)}, y_m)\}$  - training data

Classification problem:  $\gamma : \mathbf{x} \rightarrow y$

- ① Activity recognition:  $x$  - pulse rate,  $y$  - activity (walking, running, peace)
- ② Opinion mining:  $x$  - sentence,  $y$  - sentiment (positive, negative)
- ③ Trading:  $x$  - stock market,  $y$  - action (sell, buy, do nothing)

## Sequence modelling II

### Sequence labelling

- ①  $\mathbf{x} = x_1, x_2, \dots, x_n, x_i \in V$  - objects
- ②  $\mathbf{y} = y_1, y_2, \dots, y_n, y_i \in \{1, \dots, L\}$  - labels
- ③  $\{(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), (\mathbf{x}^{(2)}, \mathbf{y}^{(2)}), \dots, (\mathbf{x}^{(m)}, \mathbf{y}^{(m)})\}$  - training data
- ④ exponential number of possible solutions : if  $\text{length}(\mathbf{x}) = n$ , there are  $L^n$  possible solutions

Classification problem:  $\gamma : \mathbf{x} \rightarrow \mathbf{y}$

- ① Part of speech tagging:  $x$  - word,  $y$  - part of speech (verb, noun, etc.)
- ② Genome annotation:  $x$  - DNA,  $y$  - genes
- ③ HEP tracking:  $x$  - a set of hits with backgrounds,  $y$  - hit classification

# Sequence labelling tasks

## POS tagging and Named Entity Recognition

X (words)	the	cat	sat	on	a	mat
Y (tags)	DET	NOUN	VERB	PREP	DET	NOUN

Table: POS tagging

Alex	is	going	to	Los	Angeles
B-PER	O	O	O	B-LOC	I-LOC

Table: NER (IOB2)

Alex	travels	with	Marty	A.	Rick	to	NY	city
S-PER	O	O	B-PER	I-PER	E-PER	O	B-LOC	E-LOC

Table: NER (IOBES)

## Sequence modelling III

### Sequence transduction / transformation

- ①  $\mathbf{x} = x_1, x_2, \dots, x_n, x_i \in V_{source}$  - objects
- ②  $\mathbf{y} = y_1, y_2, \dots, y_n, y_i \in V_{target}$  - objects
- ③  $\{(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), (\mathbf{x}^{(2)}, \mathbf{y}^{(2)}), \dots, (\mathbf{x}^{(m)}, \mathbf{y}^{(m)})\}$  - training data
- ④  $\mathbf{x}^{(1)}, \mathbf{y}^{(1)}$  are of different length

Transduction problem:  $\mathbf{x}_{source} \rightarrow \mathbf{y}_{target}$


- ① Machine translation:  $x$  - sentence in German,  $y$  - sentence in English
- ② Speech recognition:  $x$  - spoken language,  $y$  - text
- ③ Chat bots:  $x$  - question,  $y$  - answer

# Traditional ML approaches to sequence modeling

- Hidden Markov Models (HMM)
- Conditional Random Fields (CRF)
- Local classifier: for each  $x$  define features, based on  $x_{-1}$ ,  $x_{+1}$ , etc, and perform classification  $n$  times

Problems:

- 1 Markov assumption: fixed length history
- 2 Computation complexity




# DL approaches to sequence modeling

- Neural networks
- Recurrent neural network and its modifications: LSTM, GRU, Highway
- 2D Convolutional Neural Network
- Transformer
- Pointer network

Problems:

- 1 Training time
- 2 Amount of training data



# Neural language model





Figure: Neural language model



## Recurrent neural network

- Input: sequence of vectors
- $x_{1:n} = x_1, x_2, \dots, x_n, x_i \in \mathbb{R}^{d_{in}}$
- Output: a single vector  
 $y_n = RNN(x_{1:n}), y_n \in \mathbb{R}^{d_{out}}$
- For each prefix  $x_{1:j}$  define an output vector  $y_j$ :  
 $y_j = RNN(x_{1:j})$
- $RNN^*$  is a function returning this sequence for input sequence  $x_{1:n}$ :  
 $y_{1:n} = RNN^*(x_{1:n}), y_i \in \mathbb{R}^{d_{out}}$



## Sequence modelling with RNN

### ① Sequence classification

Put a dense layer on top of RNN to predict the desired class of the sequence after the whole sequence is processed


$$p(l_j | \mathbf{x}_{1:n}) = \text{softmax}(RNN(\mathbf{x}_{1:n}) \times W + b)_{[j]}$$

### ② Sequence labelling

Produce an output  $y_i$  for each input RNN reads in. Put a dense layer on top of each output to predict the desired class of the input

$$p(l_j | \mathbf{x}_j) = \text{softmax}(RNN(\mathbf{x}_{1:j}) \times W + b)_{[j]}$$


## RNN unrolled



$$\begin{aligned} s_4 &= R(s_3, x_4) = R(R(s_2, x_3), x_4) = R(R(R(s_1, x_2), x_3), x_4) = \\ &= R(R(R(R(s_0, x_1), x_2), x_3), x_4) \end{aligned}$$

## Bidirectional RNN (Bi-RNN)


The input sequence can be read from left to right and from right to left.  
Which direction is better?



$$biRNN(x_{1:n}, i) = y_i = [RNN^f(x_{1:i}); RNN^r(x_{n:i})]$$

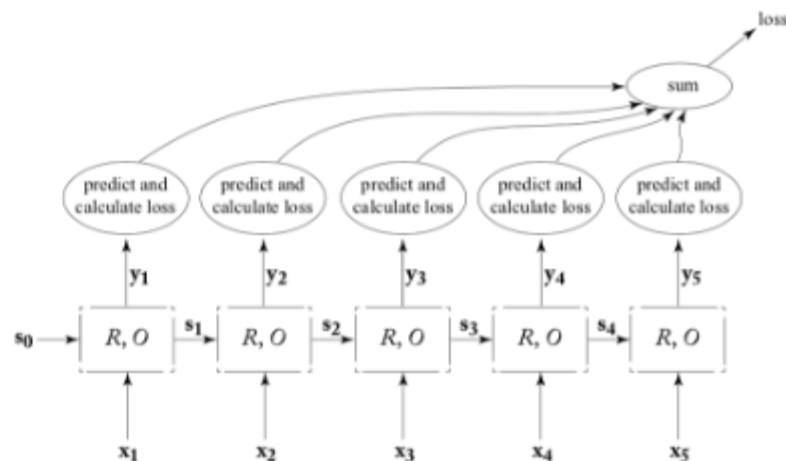
## Sequence classification

- $\hat{y}_n = O(s_n)$
- $\text{prediction} = \text{MLP}(\hat{y}_n)$
- Loss:  $L(\hat{y}_n, y_n)$
- $L$  can take any form: cross entropy, hinge, margin, etc.




## Sequence labelling

- Output  $\hat{t}_i$  for each input  $x_{1,i}$
- Local loss:  $L_{local}(\hat{t}_i, t_i)$
- Global loss:  
$$L(\hat{t}_n, t_n) = \sum_i L_{local}(\hat{t}_i, t_i)$$
- $L$  can take any form: cross entropy, hinge, margin, etc.



## Backpropagation through time



$$s_i = R(x_i, s_{i-1}) = g(s_{i-1}W^s + x_iW^x + b)$$

$$\text{Chain rule: } \frac{\partial L}{\partial w} = \frac{\partial L}{\partial p(\hat{y}_5)} \frac{\partial p(\hat{y}_5)}{\partial s_4} \left( \frac{\partial s_4}{\partial w} + \frac{\partial s_4}{\partial s_3} \frac{\partial s_3}{\partial w} + \frac{\partial s_4}{\partial s_3} \frac{\partial s_3}{\partial s_2} \frac{\partial s_2}{\partial w} + \dots \right)$$

## Vanishing gradient problem

Chain rule:  $\frac{\partial L}{\partial w} = \frac{\partial L}{\partial p(\hat{y}_5)} \frac{\partial p(\hat{y}_5)}{\partial s_4} \left( \frac{\partial s_4}{\partial w} + \frac{\partial s_4}{\partial s_3} \frac{\partial s_3}{\partial w} + \frac{\partial s_4}{\partial s_3} \frac{\partial s_3}{\partial s_2} \frac{\partial s_2}{\partial w} + \dots \right)$   
 $g$  – sigmoid

- ① Many sigmoids near 0 and 1
  - ▶ Gradients  $\rightarrow 0$
  - ▶ Not training for long term dependencies
- ② Many sigmoids  $> 1$ 
  - ▶ Gradients  $\rightarrow +\infty$
  - ▶ Not training again


Solution: gated architectures (LSTM and GRU)


## Controlled memory access

- Entire memory vector is changed:  $s_{i+1} = R(x_i, s_i)$
- Controlled memory access:  $s_{i+1} = g \odot R(x_i, s_i) + (1 - g)s_i$   
 $g \in [0, 1]^d, s, x \in \mathbb{R}^d$
- Differential gates:  $\sigma(g), g' \in \mathbb{R}^d$
- This controllable gating mechanism is the basis of the LSTM and the GRU architectures




## Long short term memory






$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$




$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$




$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$



$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

## Gated recurrent unit



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

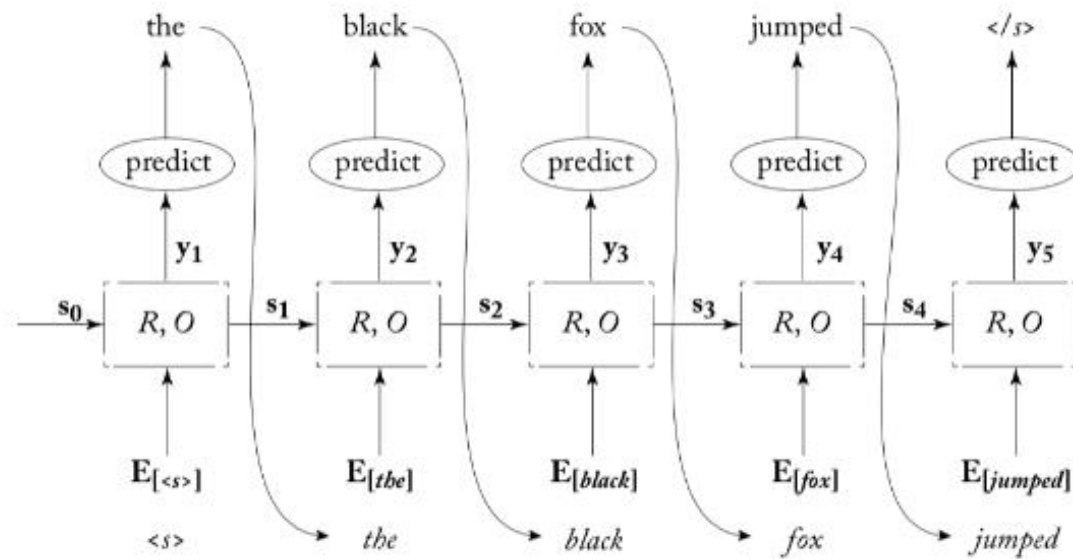
$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

## Sequence generation

Teacher forcing:  $x := \langle s \rangle x, y := x \langle /s \rangle$

$x : \langle s \rangle x_1 x_2 \dots x_n$

$y : x_1 x_2 \dots x_n \langle /s \rangle$



# Pros and cons of RNNs

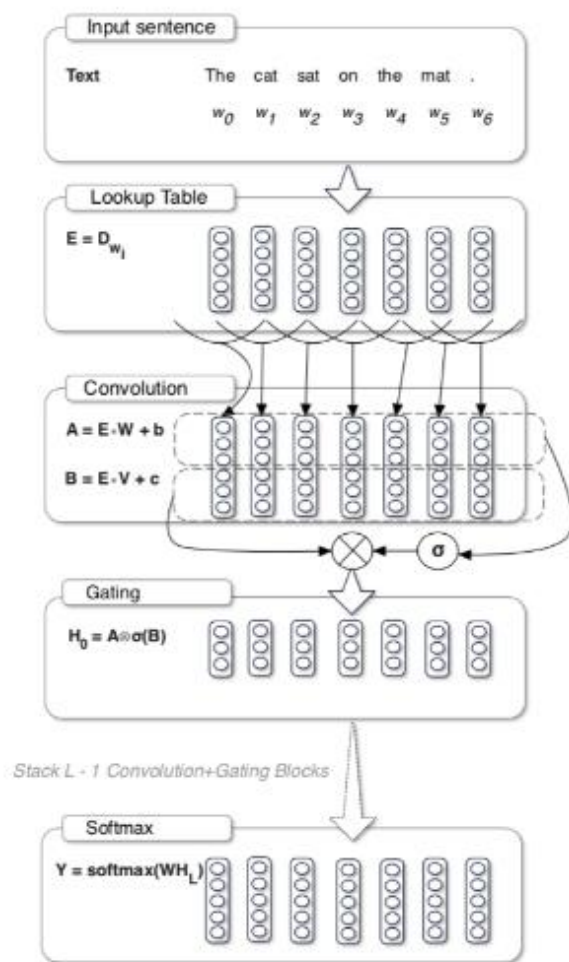
## ① Advantages:

- ▶ RNNs are popular and successful for variable-length sequences
- ▶ The gating models such as LSTM are suited for long-range error propagation

## ② Problems:

- ▶ The sequentiality prohibits parallelization within instances
- ▶ Long-range dependencies still tricky, despite gating

# Language Modeling with Gated Convolutional Networks



- **Embeddings**  $\in D^{|V| \times e}$

- **Input:**  $w_0, \dots, w_n \rightarrow E = [D_{w_0}, \dots, D_{w_n}]$

- **Hidden layers:**  $h_0, \dots, h_n$ :

$$h_l(X) = (X \times W + b) \circ \sigma(X \times V + c)$$

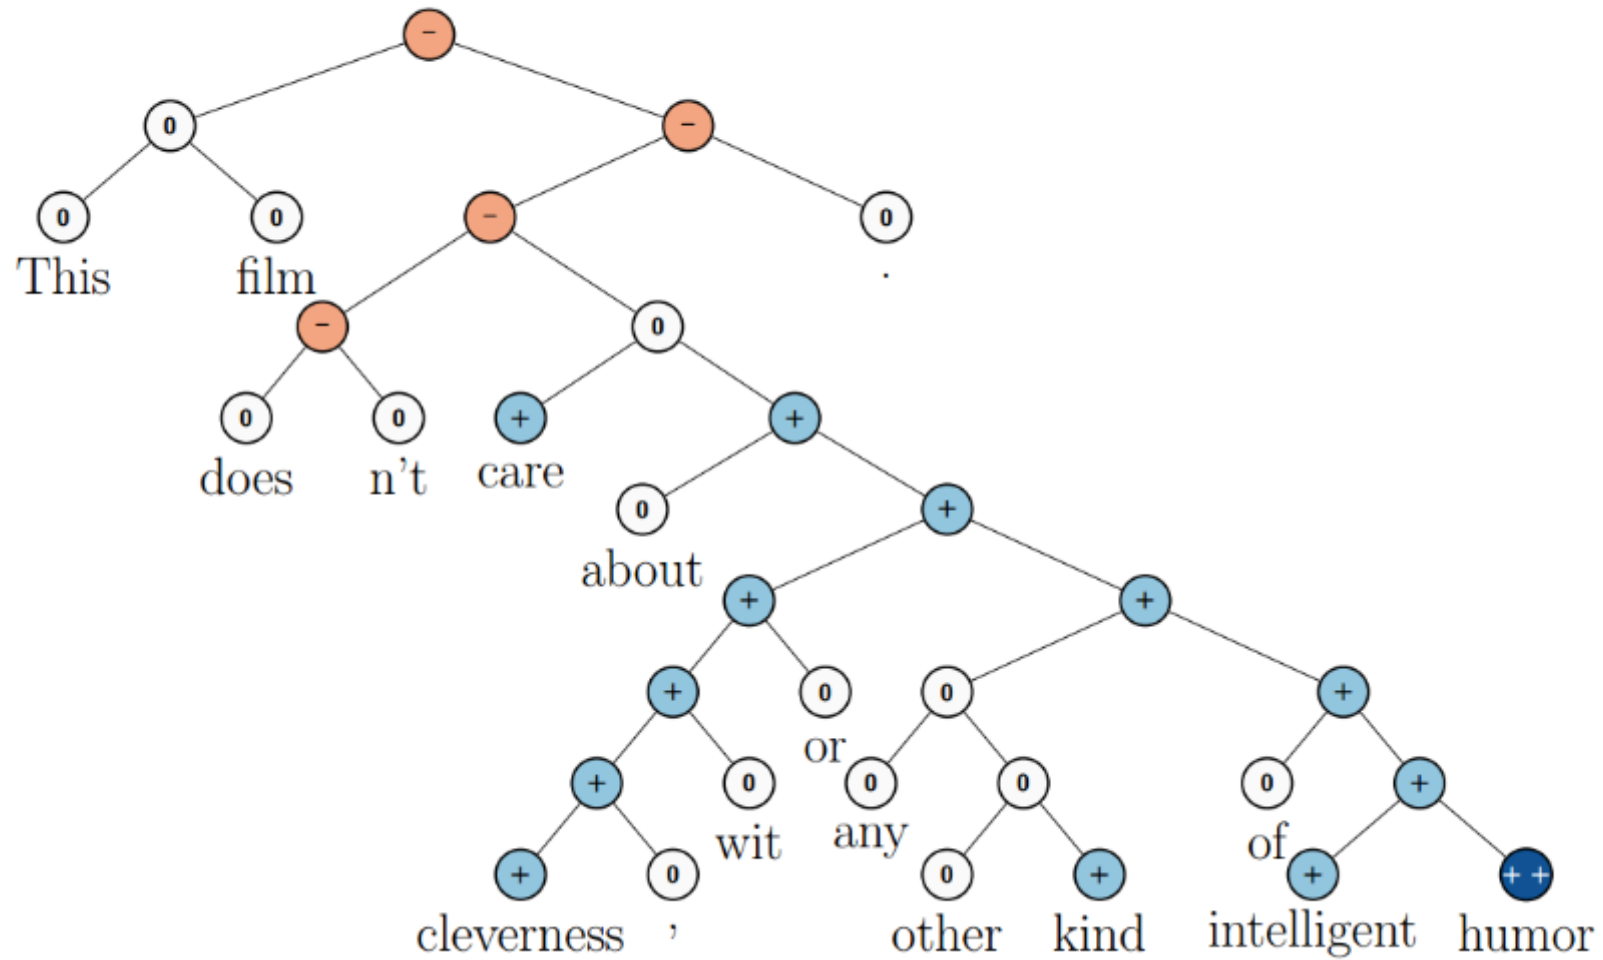
- **Gated linear unit:**  $X \circ \sigma(X)$

- **Output:**  $Y = \text{softmax}(WH_L)$

## Modeling trees with Recursive NN

- Input:  $x_1, x_2, \dots, x_n$
- A binary tree  $T$  can be represented as a unique set of triplets  $(i, k, j)$ , s.t.  $i < k < j$ ,  $x_{i:j}$  is parent of  $x_{i:k}, x_{k+1:j}$
- RecNN takes as an input a binary tree and returns as output a corresponding set of inside state vectors  $\mathbf{s}_{i:j}^A \in \mathbb{R}^d$
- Each state vector  $\mathbf{s}_{i:j}^A$  represents the corresponding tree node  $q_{i:j}^A$  and encodes the entire structure rooted at that node

## RecNN






# The Transformer

An alternative architecture to RNN which allows of parallel and faster training

- Several layers of identical modules
- Each module consists of Multi-Head Attention and Feed Forward layers
- Input: embeddings. To get embeddings for numerical input, apply any dense layer
- Positional embeddings to make use of the order of the sequence




## Scaled Dot-Product Attention


An attention function can be described as mapping a query and a set of key-value pairs to an output, where the query, keys, values, and output are all vectors:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V,$$


where the input consists of queries  $Q$  and keys  $K$  of dimension  $d_k$  and values  $V$  of dimension  $d_v$



# Modern RNNs: RetNet



(a) Parallel representation.



(b) Recurrent representation.




Figure 3: Dual form of RetNet. “GN” is short for GroupNorm.

The **RetNet model** (short for **Retention Network**, introduced by Microsoft Research in 2023) is a new sequence modeling architecture proposed as an alternative to Transformers. Its main idea is to keep the parallelizability of Transformers while introducing recurrent-like efficiency for very long sequences.

# Modern RNNs: RetNet

## 1. Retention Mechanism

1. RetNet replaces the Transformer's self-attention with a new operator called **retention**.
2. It computes token representations with a **decay-weighted recurrence**, which combines elements of attention (parallelizable) and RNNs (sequential memory).
3. This allows it to **remember long-term dependencies** without quadratic complexity.

## 2. Linear Time & Memory Complexity

1. Like linear-attention methods, RetNet has  **$O(n)$**  complexity in sequence length (vs.  $O(n^2)$  for standard Transformers).
2. Scales well for very long contexts (tens of thousands of tokens).

## 3. Parallel + Recurrent Modes

1. RetNet can run in **parallel mode** (fast training, like Transformers) and **recurrent mode** (efficient autoregressive inference, like RNNs).
2. This duality makes it especially practical for **deployment**.

# Modern RNNs: Mamba

## 1. State Space Model (SSM) Backbone

- Mamba is based on **selective state space models (SSMs)** rather than attention.
- SSMs update a hidden state with each token, similar to RNNs, but with a **structured linear dynamical system** that allows efficient sequence modeling.

## 2. Selectivity Mechanism

- Introduces a **data-dependent gating/selectivity** mechanism, letting the model decide which past information to keep or forget.
- This makes memory usage **adaptive** rather than uniform (unlike RetNet's exponential decay).

## 3. Linear Time Complexity

- Like RetNet, Mamba achieves  **$O(n)$**  complexity in sequence length.
- Memory and compute scale linearly, making it suitable for **very long contexts (hundreds of thousands of tokens)**.

## 4. Hardware Efficiency


- The architecture is designed to map efficiently to GPUs/TPUs.
- Thanks to structured kernels, Mamba is **much faster than Transformers** for long sequences, both in training and inference.

## 5. Continuous State Representation


- Maintains a **continuous hidden state** (like an RNN) that evolves over time.
- This allows **constant-time per-token inference**, unlike attention models which need growing context windows.

# Modern RNNs: Mamba

## Copying



## Selective Copying



## Induction Heads



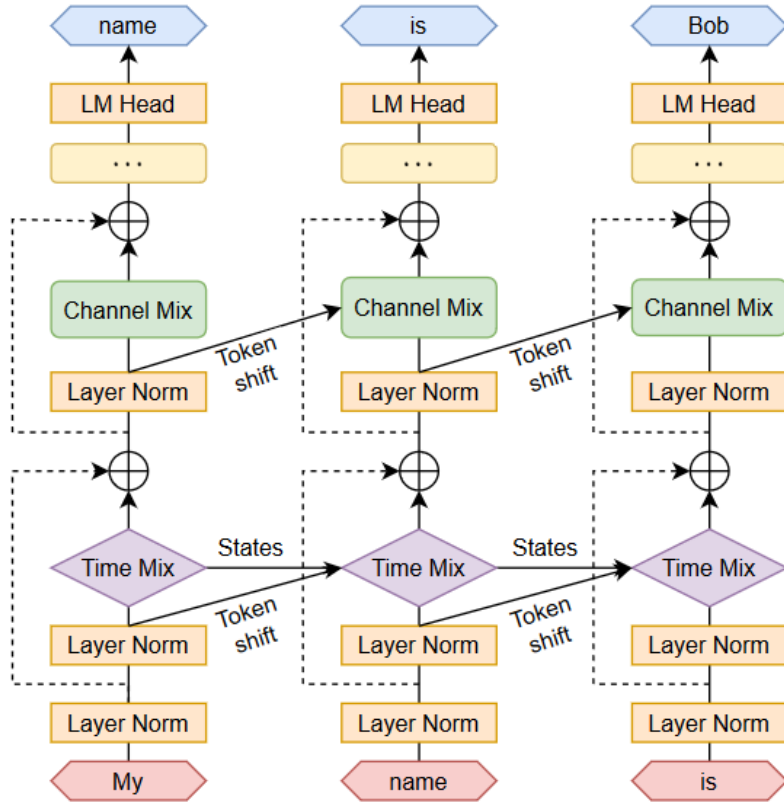


Figure 2: (Left) The standard version of the Copying task involves constant spacing between input and output elements and is easily solved by time-invariant models such as linear recurrences and global convolutions. (Right Top) The Selective Copying task has random spacing in between inputs and requires time-varying models that can *selectively* remember or ignore inputs depending on their content. (Right Bottom) The Induction Heads task is an example of associative recall that requires retrieving an answer based on context, a key ability for LLMs.

# Modern RNNs: Mamba



# Modern RNNs: RWKV



- $R$ : The **Receptance** vector acts as the receiver of past information.
- $W$ : The **Weight** signifies the positional weight decay vector, a trainable parameter within the model.
- $K$ : The **Key** vector performs a role analogous to  $K$  in traditional attention mechanisms.
- $V$ : The **Value** vector functions similarly to  $V$  in conventional attention processes.

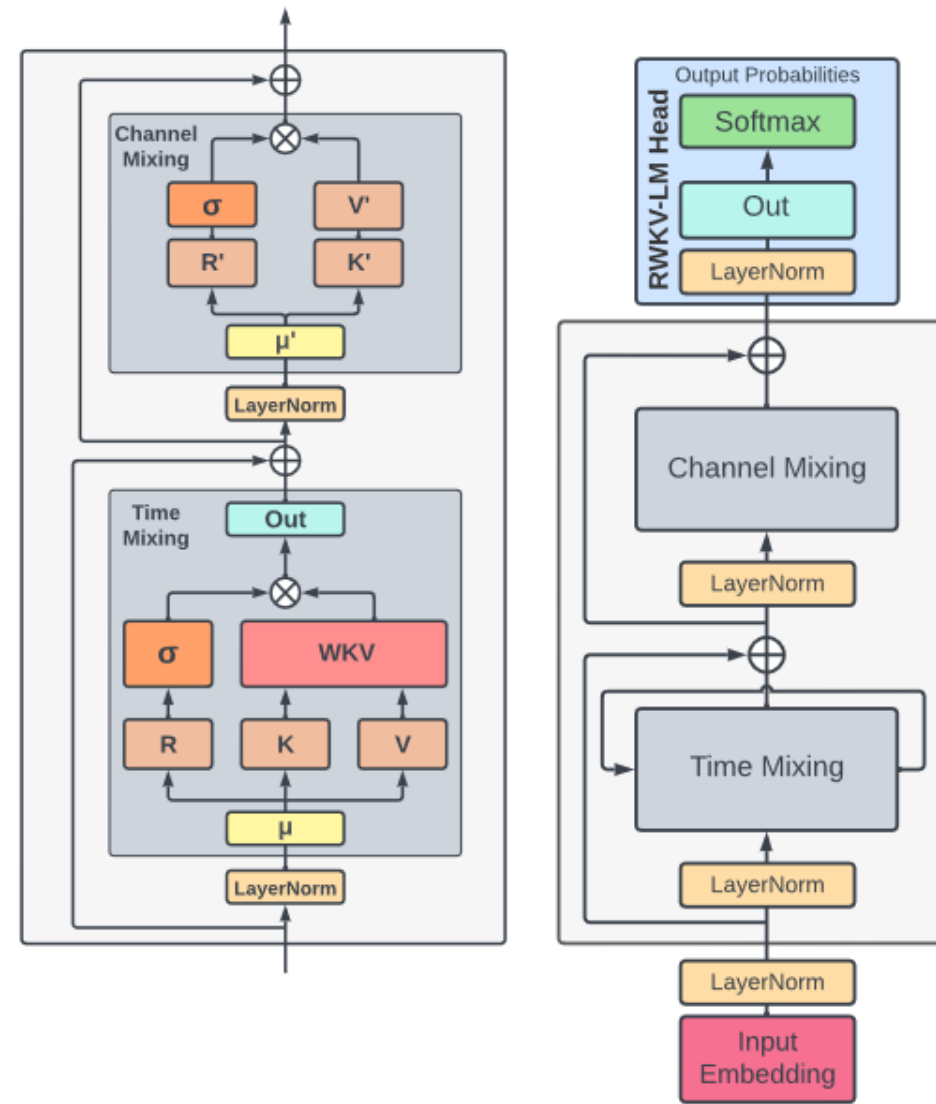


Figure 2: Elements within an RWKV block (left) and the complete RWKV residual block, equipped with a final head for language modeling (right).



# Modern RNNs: RWKV

## 1. RNN–Transformer Hybrid

1. RWKV is built as a **recurrent neural network** but trained with **Transformer-like weight sharing and architecture patterns**.
2. It can process sequences one token at a time (like an RNN) while keeping some of the expressivity of attention models.

## 2. Linear Complexity

1. Like RetNet and Mamba, RWKV has  **$O(n)$**  time and memory complexity.
2. It avoids quadratic attention costs, making it efficient for **very long sequences**.

## 3. Time-Mixing & Channel-Mixing Blocks

1. Instead of self-attention, RWKV uses **time-mixing** (temporal recurrence across tokens) and **channel-mixing** (feedforward across hidden dimensions).
2. This plays a role similar to attention + MLP in Transformers, but in a recurrent-friendly way.

## 4. Recurrent Inference Mode

1. RWKV can run **autoregressively like an RNN**: constant memory per token, only needing the hidden state.
2. This makes inference **highly efficient** for deployment on CPUs and edge devices.

## 5. Transformer-Like Training

1. Despite its RNN nature, RWKV can be trained with **parallelism similar to Transformers** using special weight formulations.
2. This enables large-scale training with GPUs/TPUs.

# Modern RNNs: ARMT

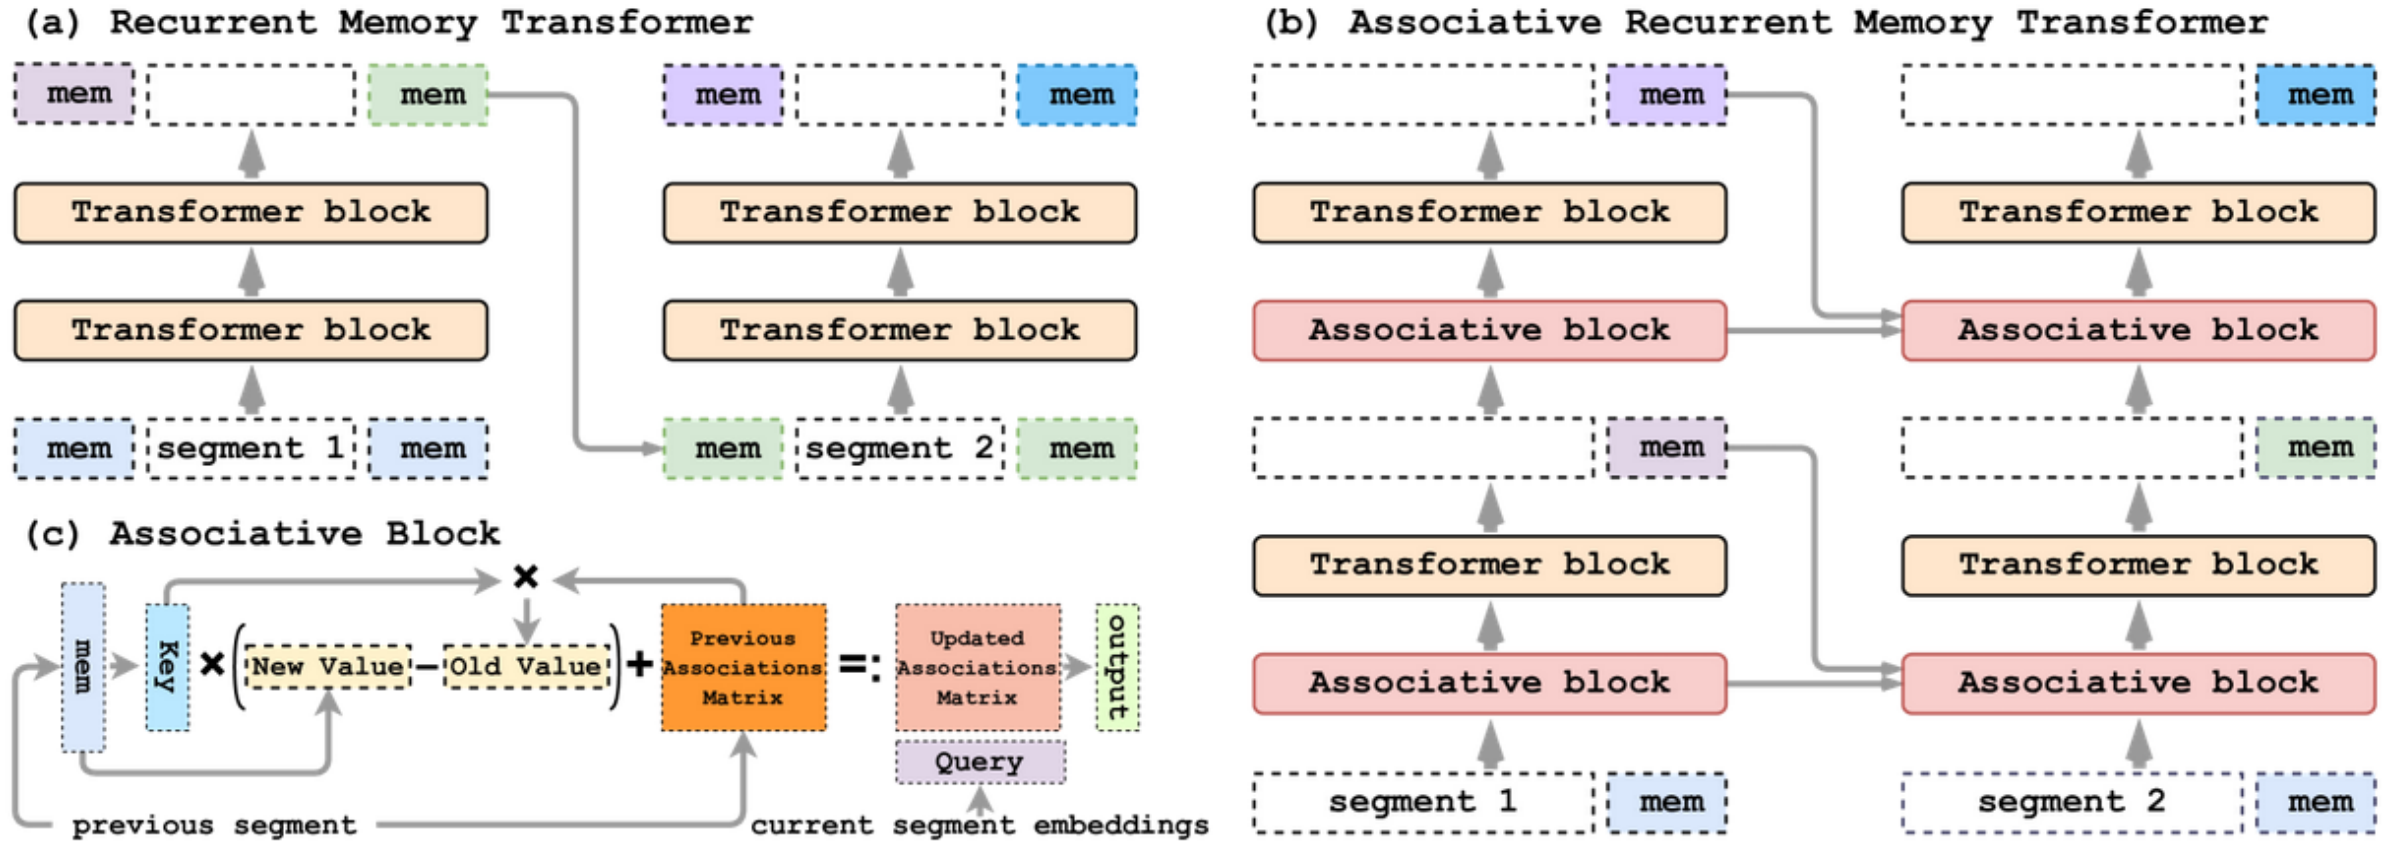


Figure 1: **ARMT augments the transformer's layers with associative memory.** (a) RMT architecture. (b) ARMT adds associative memory processing to each layer. (c) Associative memory is updated with layerwise memory representations.

**Associative Recurrent Memory Transformer (ARMT)** — 2024. Hybrid model combining local Transformer self-attention + **segment-level recurrence** to store and retrieve information across large contexts. Designed to handle very long sequences. “**Constant time for processing new information at each time step**” is claimed.